

# PARALLEL & DISTRIBUTED GRAPH PROCESSING

Analysis of Large Scale Social Networks

Bart Thijs

The focus of this lecture is on the different approaches taken to implement a distributed graph processing framework.

Pregel: A system for large-scale graph processing

Overview of Distributed Graph Processing (Kalavri et al., 2018)

The implementation of a graph algorithm in a DGP (Thijs, 2017)


GraphX as a DGP-framework.

# DISTRIBUTED GRAPH PROCESSING

Several white lines of varying lengths and slopes are drawn in the bottom right corner of the slide, creating a modern, abstract graphic element.

## Pagerank :

- ▶ Rank webpages based on number and quality of links
- ▶ More important pages receive more (important) links
- ▶ Value propagation algorithm
- ▶ Evenly distributed outgoing links
- ▶ Random clicking  $(1-d)=15\%$

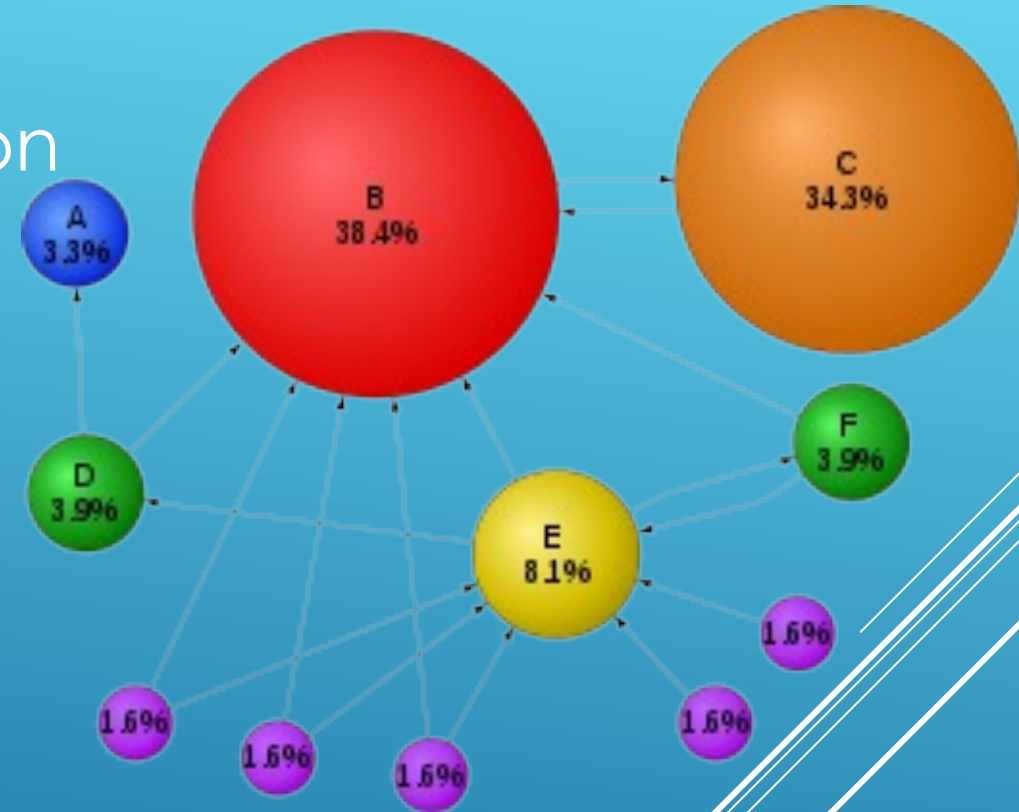


$$PR(p_i) = \frac{1-d}{N} + d \sum_{p_j \in M(p_i)} \frac{PR(p_j)}{L(p_j)}$$

FIRST

## Pagerank :

- ▶ Matrix Approach starts from transition matrix
- ▶ In fact, the Pagerank vector is the eigenvector associated with the eigenvalue of 1.



$$PR(p_i) = \frac{1-d}{N} + d \sum_{p_j \in M(p_i)} \frac{PR(p_j)}{L(p_j)}$$

FIRST

# How can we efficiently calculate Pagerank?

- ▶ Huge data set, too big for regular hardware / software
- ▶ Dynamic: ever growing and changing

Need for specialized graph processing systems vs algebraic approach:

- ▶ Express program logic as close as possible to data model

## WHY?

## Three challenges for graph analytical frameworks

1. Capacity (Storage: Memory and Disk)
2. Performance (due to RAM-reads exceeding cache)
3. Implementation (parallelization, small tasks at node and edge level)

THREE CHALLENGES (HONG ET AL)

A series of white diagonal lines of varying lengths and thicknesses, located in the bottom right corner of the slide.

2 possible approaches for parallelization of the graph analysis:

- ▶ Multi-core computers with shared memory
  - ▶ HPC (Cerebro @ KU Leuven),
- ▶ Distributed computer systems
  - ▶ Hadoop/Spark (EMR @ Amazon)

SOLUTION?

A series of white diagonal lines of varying lengths and thicknesses, located in the bottom right corner of the slide.

Implementing graph algorithms becomes very difficult:

- ▶ Right level or form of parallelization or distribution with appropriate communication and synchronyzation
- ▶ Correct implementation of described algorithm
- ▶ Optimized for each particular environment
- ▶ Fault tolerance

Who has all these skills/expertise?

SOLUTION?



## Shared memory frameworks

- SNAP
- GraphLab
- Ligra
- Green Marl

# PARALLEL GRAPH PROCESSING

Several white lines of varying lengths and slopes are positioned in the bottom right corner of the slide, creating a modern, abstract graphic element.

## Frameworks for distributed graph processing

- ▶ High level of abstraction
- ▶ Partitioned view of data and methods for read/write in and communication between partitions.

Examples:

- Pregel (Google)
- Giraph (Facebook)
- GraphX (Spark)

# DISTRIBUTED GRAPH PROCESSING

Several white lines of varying lengths and slopes are drawn in the bottom right corner of the slide, creating a modern, abstract graphic element.

## Definition

- ▶ A **distributed graph programming** model is an abstraction of the underlying computing infrastructure that allows for the definition of graph data structures and the expression of graph algorithms
  - ▶ Data partitioning and communication is hidden

KALAVRI (2018)

A series of white diagonal lines of varying lengths and thicknesses are positioned in the bottom right corner of the slide, creating a modern, abstract graphic element.

## Bulk-Synchronous Parallel (Vaillant, 1990)

Combination three attributes:

- A. A number of components/instances performing processing or memory functions:
  - 1. perform local computational steps:
  - 2. Transmit messages to other components:
  - 3. Receive messages as input for next superstep
- B. A router that delivers messages point to point between components
- C. Facilities for synchronizing all components at regular intervals through the use of supersteps. (Not required)

A superstep is finalized after all components finish their task

## DISTRIBUTED GRAPH PROCESSING

Large scale graph processing at Google:

The BSP concept is ported towards networked data:

A superstep consists of tasks performed at the vertices

1. Input of messages from previous superstep to nodes
2. Start execution at node-level of computations and change the state of vertex and (outgoing) edges
3. Collect messages to nodes (along (outgoing) edges).

Vertices are either in Active or Inactive State

# DISTRIBUTED GRAPH PROCESSING: PREGEL

```
class PageRankVertex
  : public Vertex<double, void, double> {
public:
  virtual void Compute(MessageIterator* msgs) {
    if (superstep() >= 1) {
      double sum = 0;
      for (; !msgs->Done(); msgs->Next())
        sum += msgs->Value();
      *MutableValue() =
        0.15 / NumVertices() + 0.85 * sum;
    }

    if (superstep() < 30) {
      const int64 n = GetOutEdgeIterator().size();
      SendMessageToAllNeighbors(GetValue() / n);
    } else {
      VoteToHalt();
    }
  }
};
```

**Figure 4: PageRank implemented in Pregel.**

```

class ShortestPathVertex
: public Vertex<int, int, int> {
void Compute(MessageIterator* msgs) {
    int mindist = IsSource(vertex_id()) ? 0 : INF;
    for (; !msgs->Done(); msgs->Next())
        mindist = min(mindist, msgs->Value());
    if (mindist < GetValue()) {
        *MutableValue() = mindist;
        OutEdgeIterator iter = GetOutEdgeIterator();
        for (; !iter.Done(); iter.Next())
            SendMessageTo(iter.Target(),
                          mindist + iter.GetValue());
    }
    VoteToHalt();
}
};

```

**Figure 5: Single-source shortest paths**

```

class MinIntCombiner : public Combiner<int> {
    virtual void Combine(MessageIterator* msgs) {
        int mindist = INF;
        for (; !msgs->Done(); msgs->Next())
            mindist = min(mindist, msgs->Value());
        Output("combined_source", mindist);
    }
};

```

**Figure 6: Combiner that takes minimum of message values.**

## Six models

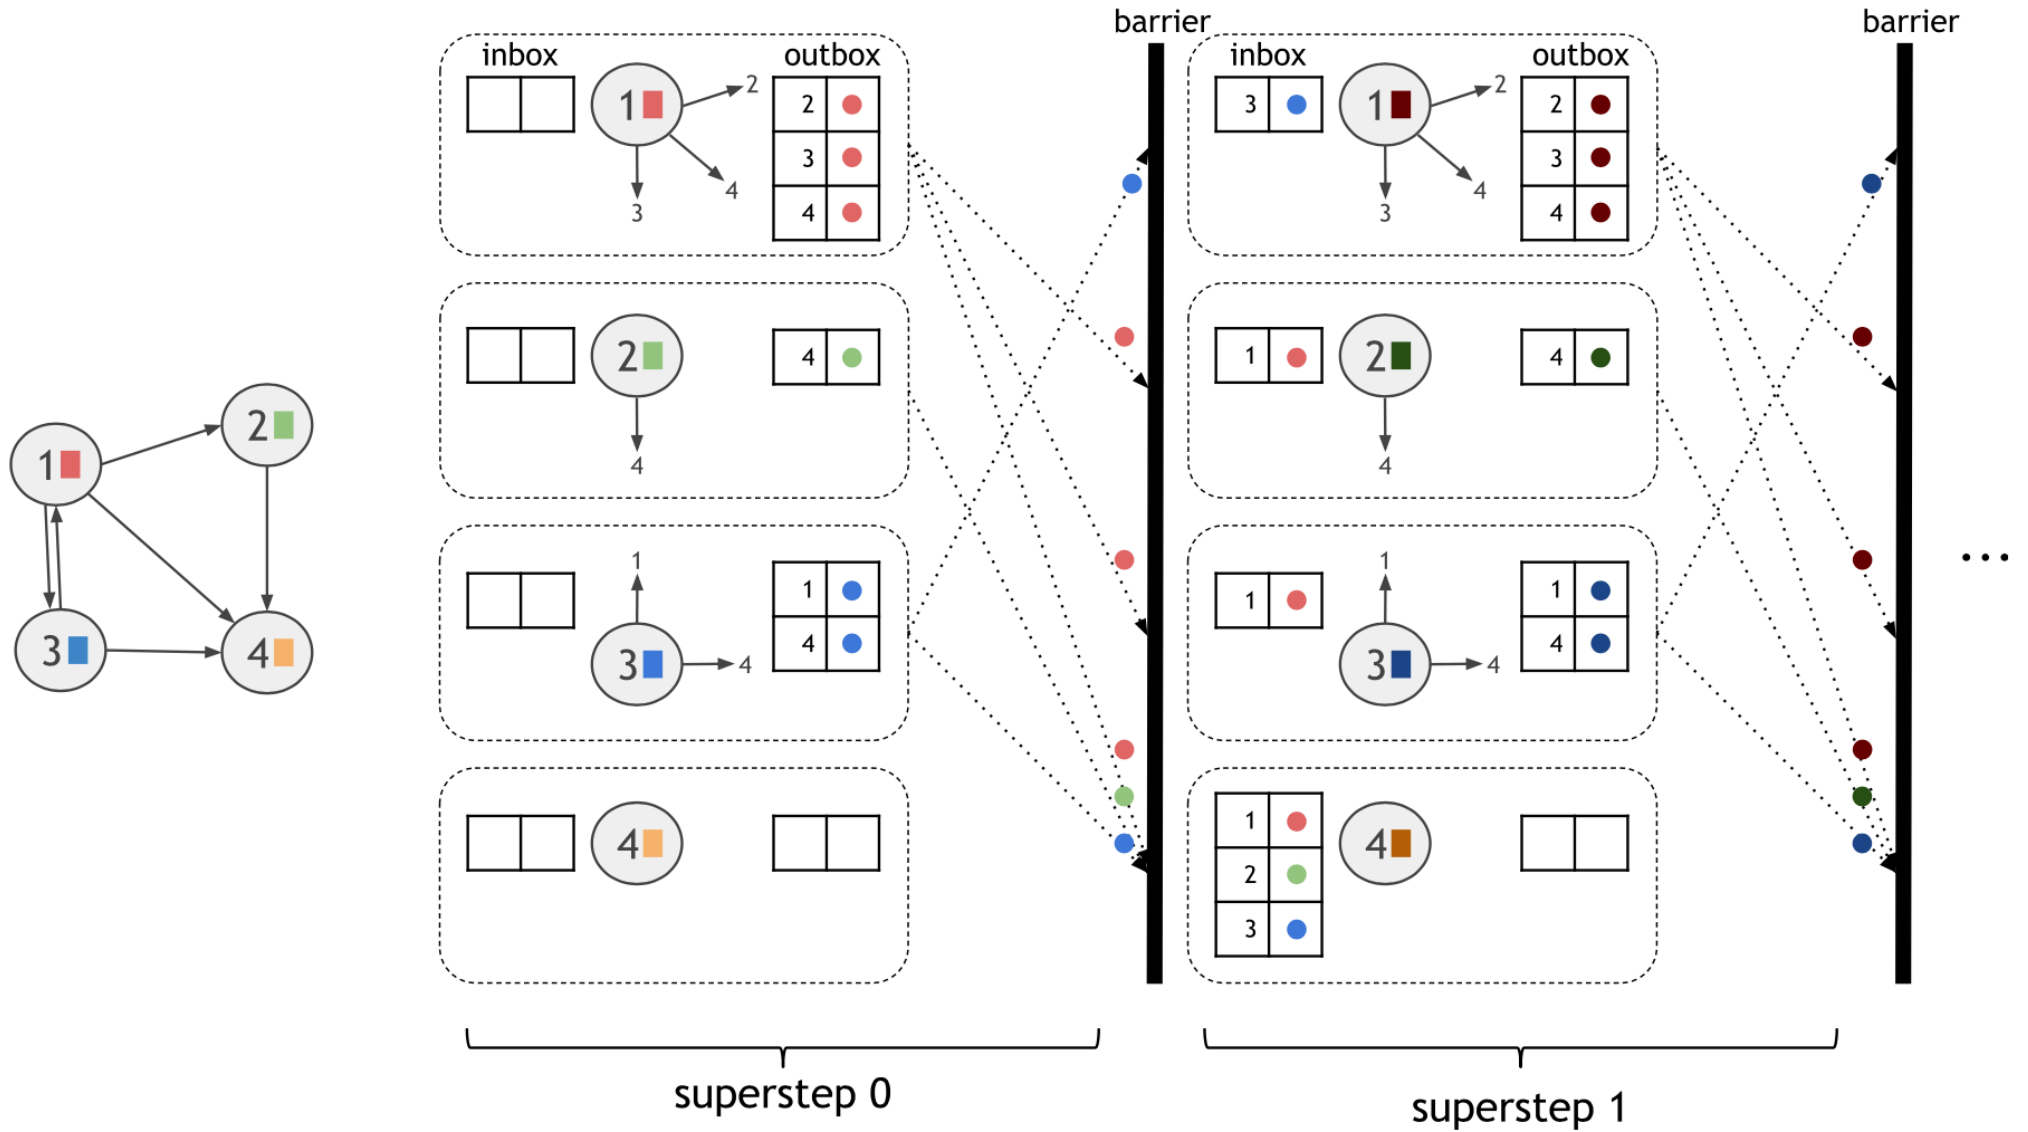
- ▶ Vertex Centric
- ▶ Scatter-Gather
- ▶ Gather-Sum-Apply-Scatter
- ▶ *Subgraph-Centric*
- ▶ *Filter-Process*
- ▶ *Graph Traversals*

KALAVRI (2018)

A series of white diagonal lines of varying lengths and thicknesses are positioned in the bottom right corner of the slide, creating a modern, abstract graphic element.



## ► Vertex Centric



---

## Algorithm 2. PageRank Vertex Function

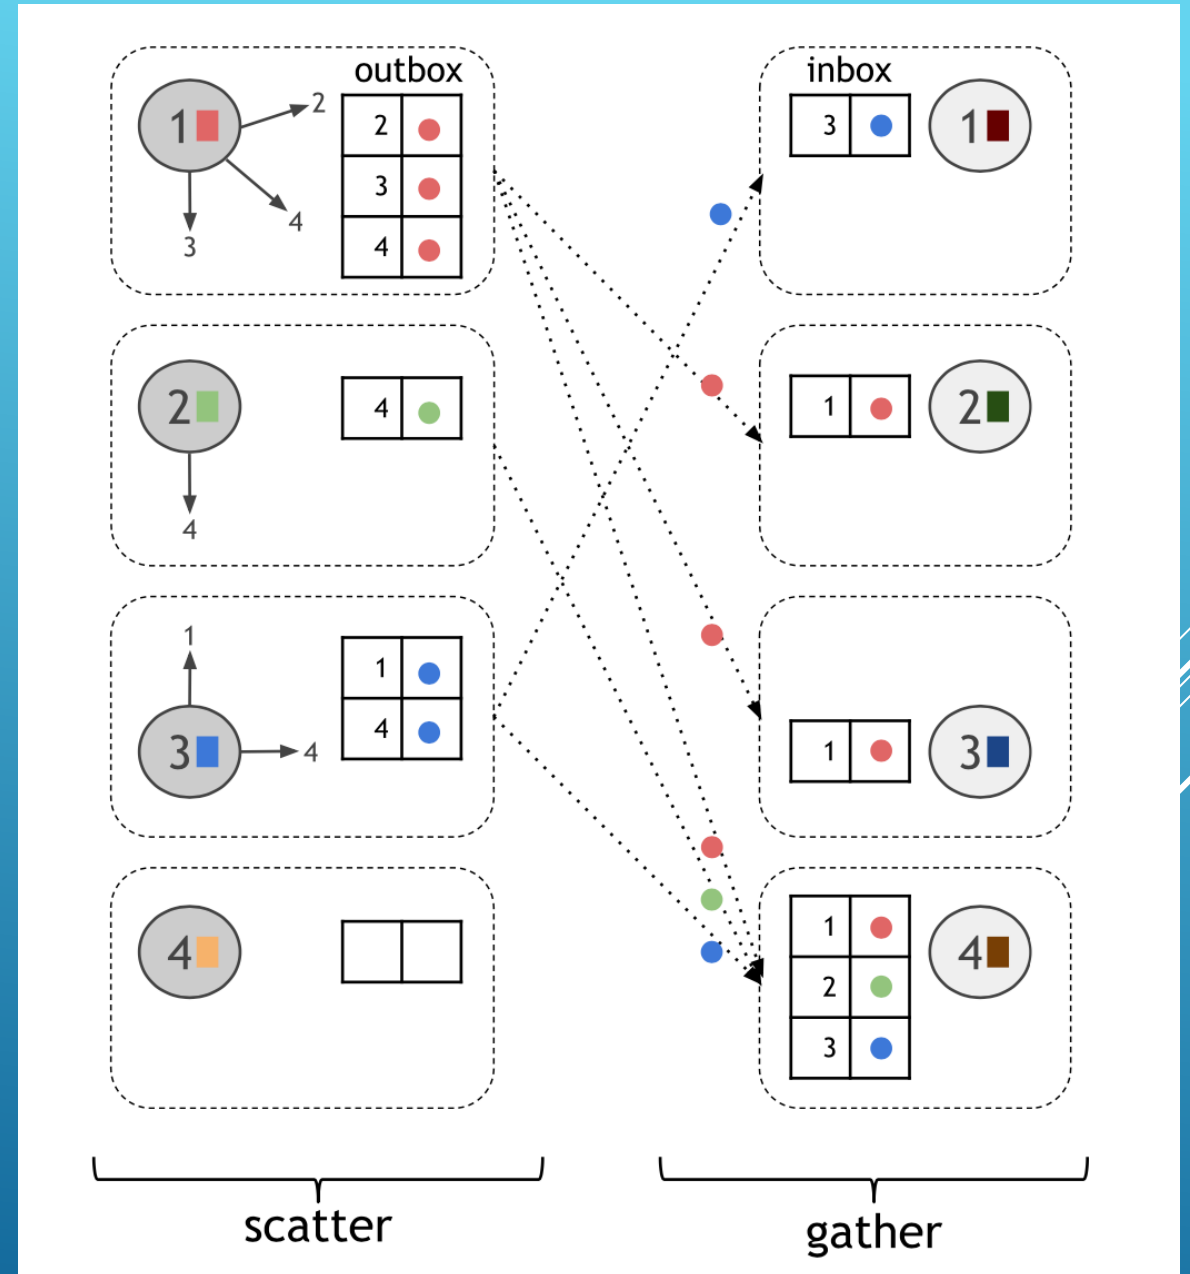
---

```
void compute(Iterator[double] messages):  
     $outE = getOutEdges().size()$   
    if  $superstep() > 0$  then  
        double  $sum = 0$   
        for  $m \in messages$  do  
             $sum \leftarrow sum + m$   
        end for  
         $setValue(0.15/numVertices() + 0.85 * sum)$   
    end if  
    if  $superstep() < 30$  then  
        for  $e \in getOutEdges()$  do  
             $sendMessageTo(e.target(), getValue()/outE)$   
        end for  
    else  
         $voteToHalt()$   
    end if
```

---

# Scatter Gather

- ▶ barrier after the Gather
- ▶ Gather has access to the individual incoming messages
- ▶ No access to in and outgoing messages.



---

## Algorithm 4. PageRank Scatter and Gather Functions

---

**void scatter():**

$outEdges = getOutEdges().size()$

**for**  $edge \in getOutEdges()$  **do**

$sendMessageTo(edge.target(), getValue()/outEdges)$

**end for**

**void gather(Iterator[double] messages):**

**if**  $superstep() < 30$  **then**

**double**  $sum = 0$

**for**  $m \in messages$  **do**

$sum \leftarrow sum + m$

**end for**

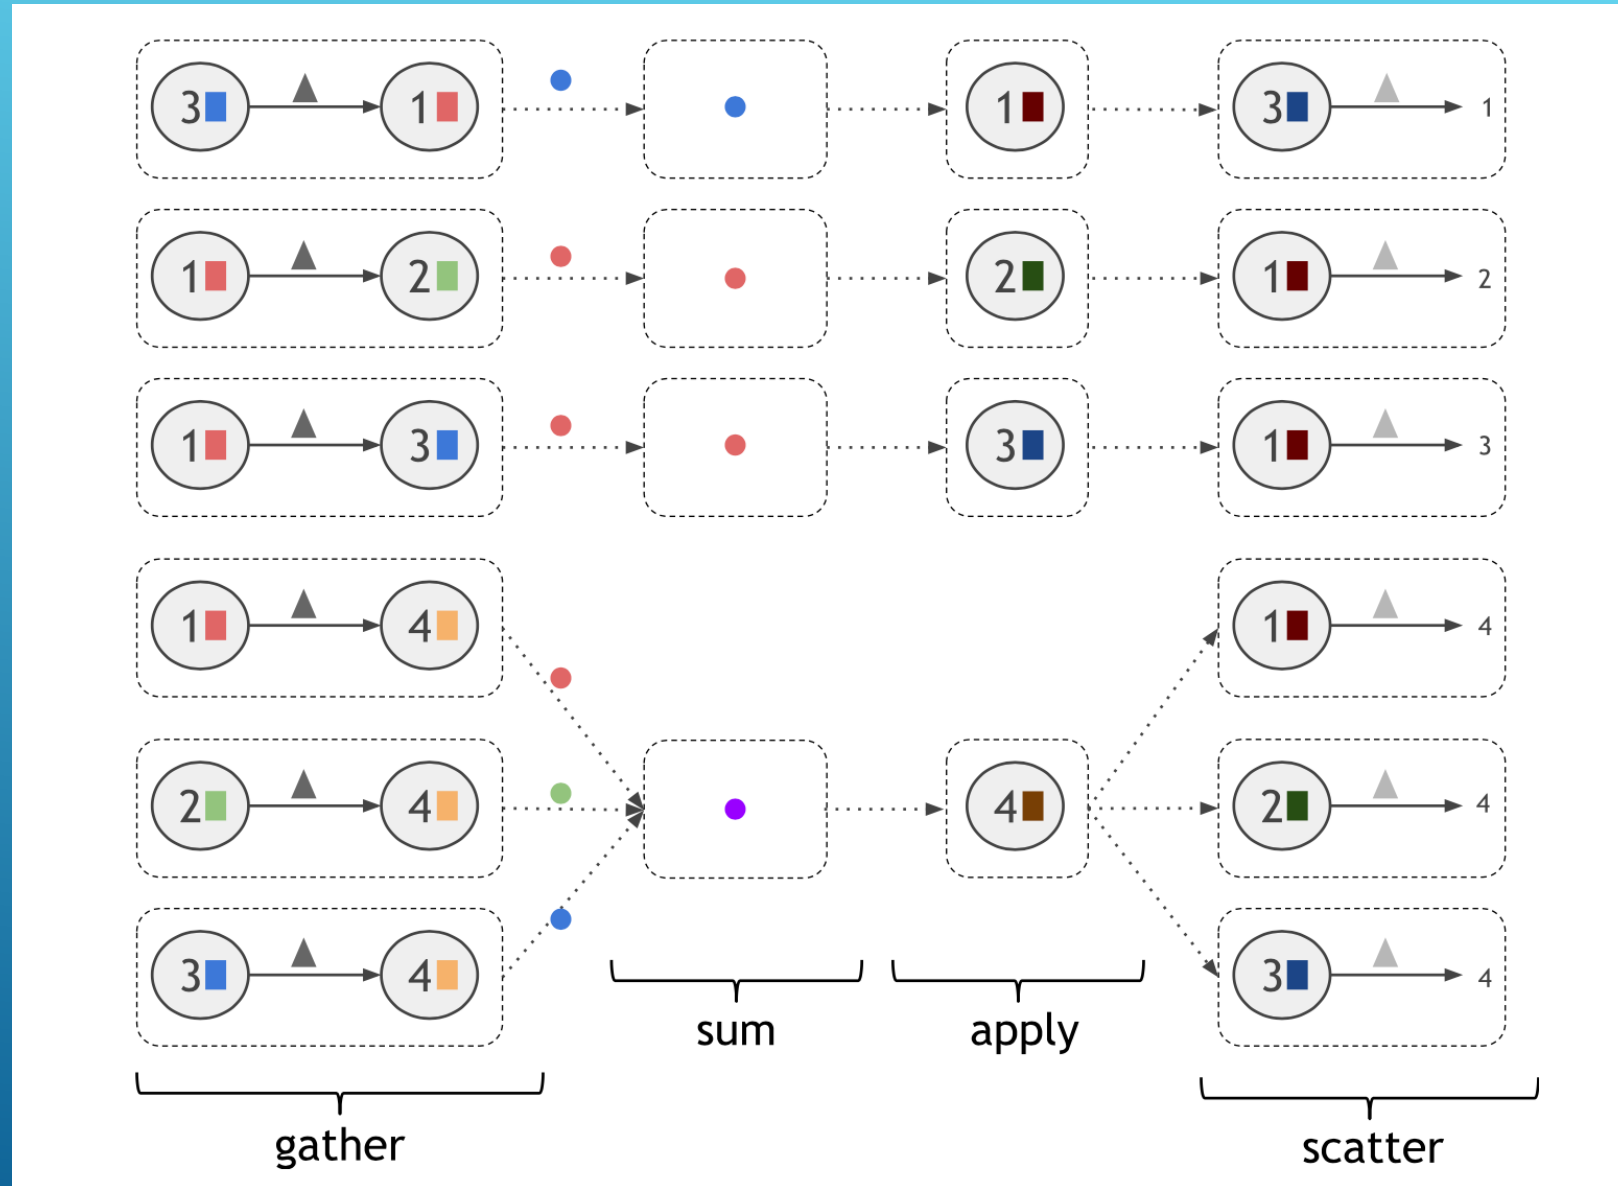
$setValue(0.15/numVertices() + 0.85 * sum)$

**end if**

---

# Gather-Sum-Apply-Scatter

- ▶ Parallelizing over the edges in the graph
- ▶ Sum = commutative & associative function
- ▶ Gather - Sum is similar to Map-Reduce



---

### Algorithm 5. GAS Model Semantics

---

**Input:** directed graph  $G=(V, E)$

$a_v \leftarrow \text{empty}$

**for**  $v \in V$  **do**

**for**  $n \in N_v^{in}$  **do**

$a_v \leftarrow \text{sum}(a_v, \text{gather}(S_v, S_{(v,n)}, S_n))$

**end for**

$S_v \leftarrow \text{apply}(S_v, a_v)$

$S_{(v,n)} \leftarrow \text{scatter}(S_v, S_{(v,n)}, S_n)$

**end for**

---

---

### Algorithm 6. PageRank Gather, Sum, Apply

---

**double gather(double src, double edge, double trg):**

**return**  $\text{trg.value} / \text{trg.outNeighbors}$

**double sum(double rank1, double rank2):**

**return**  $\text{rank1} + \text{rank2}$

**double apply(double sum, double currentRank):**

**return**  $0.15 + 0.85 * \text{sum}$

---

A nearest neighbour search algorithm, implemented within a DGP-framework (Thijs, 2017)

Bi-partite directed network connecting scientific papers with cited references.

Goal is to detect similar papers based on the number of joint references normalized by the number of references of each publication.


# DISTRIBUTED GRAPH PROCESSING

Several white lines of varying lengths and orientations are positioned in the bottom right corner of the slide, creating a modern, abstract graphic element.

Implementation in a GAS approach with 3 distinct supersteps


1. Paper sends degree and identifier to cited reference
2. Cited reference sends back list with degrees and identifiers to its citing papers
3. Papers calculate similarity with other papers and sends message to make connection.

Amount messages depend on degree



# DISTRIBUTED GRAPH PROCESSING

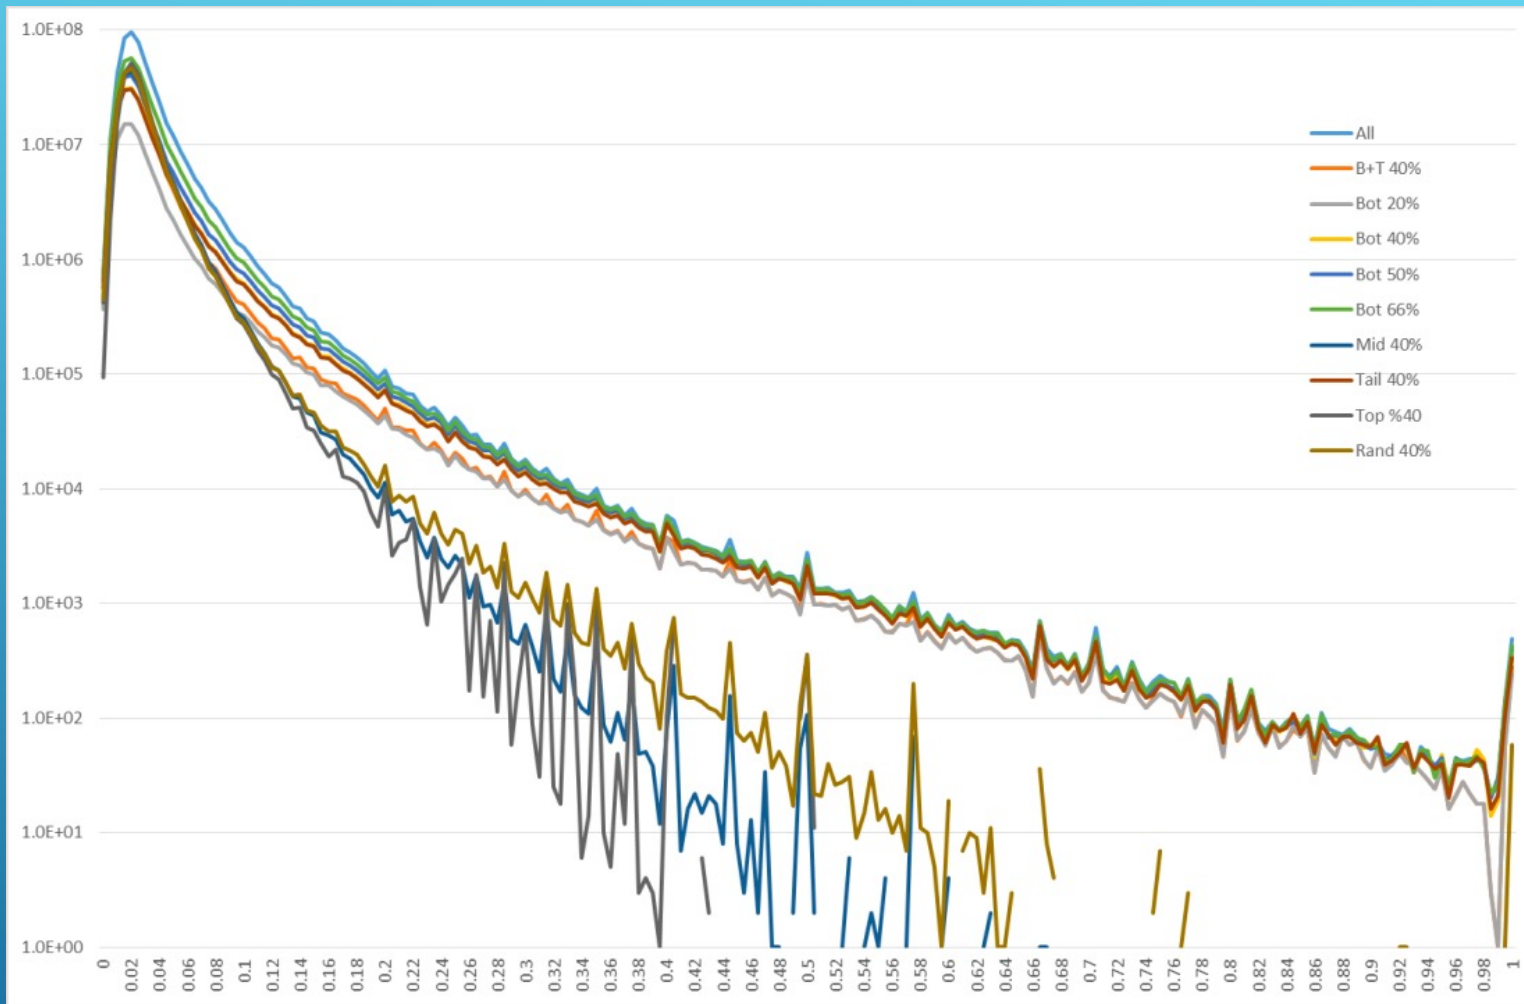




Nodes can use specific functions to reduce information in message and to select recipients of message.

High degree nodes become less likely to receive message.  
Identifiers are sorted and messages only contain list with lower ranked ID

## DISTRIBUTED GRAPH PROCESSING




# DISTRIBUTED GRAPH PROCESSING

## GraphX:

Inspired by the GAS approach and introduces for Vertex Cuts which ensures equal distribution of edges across workers.

Combiners reduce message cost across instances/workers: combiner functions must be commutative and associative.



# DISTRIBUTED GRAPH PROCESSING: GRAPHX

## Review GraphX

- ▶ Does their Pregel implementation follow the genuine Pregel model?
- ▶ How many functions are required for the implementation of the Pregel function.

DISTRIBUTED GRAPH PROCESSING

Several thin, parallel white lines are drawn diagonally across the bottom right corner of the slide, extending from the middle of the right edge towards the bottom left.

## Pregel in GraphX: Each Superstep

1. Aggregate received Messages
2. Calculate new attributes
3. Send messages to neighbours

```
def Pregel(graph: Graph[V,E],
           initialMsg: M,
           vprogf: ((Id,V), M) => V,
           sendMsgf: Edge[V,E] => Option[M],
           combinef: (M,M) => M,
           numIter: Long): Graph[V,E] = {
  // Initialize the messages to all vertices
  var msgs: RDD[(Vid, A)] =
    graph.vertices.map(v => (v.id, initialMsg))

  // Loop while there are messages
  var i = 0
  while (msgs.count > 0 && i < maxIter) {
    // Receive the message sums on each vertex
    graph = graph.updateVertices(msgs, vprogf)

    // Compute and combine new messages
    msgs = graph.aggregateNeighbors(sendMsgf,
                                     combinef)
    i = i + 1
  }
}
```

# DISTRIBUTED GRAPH PROCESSING: GRAPHX

## Page Rank in GraphX: Each Superstep

- Initial message
- 1. Combine Messages
- 2. Update Rank
- 3. Send Messages

```
// Load and initialize the graph
val graph = Graph.load('hdfs://webgraph.tsv')
var prGraph = graph.updateV(graph.degrees(OutEdges),
    (v,deg) => (v.id,(deg, 1.0)) // Initial rank=1

// Execute PageRank
prGraph = Pregel(prGraph,
    1.0, // Initial message is 1.0
    vprogf = // Update Rank
        (v, msg) => (v.deg, 0.15 + 0.85 * msg),
    sendMsgf = // Compute Msg
        e => e.src.rank/e.src.deg,
    combinef = // Combine msg
        (m1, m2) => m1 + m2,
    10) // Run 10 iterations

// Display the maximum PageRank
print(prGraph.vertices.map(v=>v.rank).max)
```

# DISTRIBUTED GRAPH PROCESSING: GRAPHX